

1. Intro
 2. Research Questions
 3. Experiments
 4. Results
 5. Answers
 6. Conclusions
-
- ```
graph TD; A[1. Intro] --> B[2. Research Questions]; B --> C[3. Experiments]; C --> D[4. Results]; D --> E[5. Answers]; E --> F[6. Conclusions]; C --> B; D --> C;
```

# PAPER OUTLINE

## 1. Intro

- a) initial plan
- b) Dynamo vs Big Table
- c) → Dynamo vs Cosmos

## 2. Research Questions

- a) Are we getting what we're paying for w.r.t provisioning & capacity
- b) Efficiency of use w.r.t. latency & throughput
- c) How easy is it to implement connections in software or other applications? How easy is it to push/pull data? (Initially & finally) How easy is it to do large updates? Streams?
- d) How do provisioning costs & implementations compare?
- e) What kinds of replication are available and how are they accomplished?

S3  
for AWS

### 3. Experiments

#### a) Types

- i) Dynamo
- ii) Cosmos w/o A.Z.  
replication
- iii) Cosmos w/ A.Z  
replication

#### b) Metrics

- i) Throughput ( $R/s$ )  $\rightarrow$  provisioning
- ii) Latency (ms)  $\rightarrow$  queries

#### c) Throughput Conditions

- i) Push
- ii) Dump

#### d) Latency Conditions

- i) no load
- ii) half push
- iii) half pull
- iv) full push
- v) full pull